

COMP 755 Advanced Operating Systems  
Fall 2025  
Philip S. Bell

## Mid-Semester Project Progress Report

What to include:

- Group information
  - Solo: Philip S. Bell
- Project title and topic:

### Decentralized Load Balancer Simulation:

Implement a distributed, routerless peer simulation in C++ where each node (thread or process) exchanges load information with peers and dynamically routes work requests to less-busy neighbors — modeling the decentralized coordination logic of my future Mixture of Experts (MoE) system for my PhD research project.

### Timeline:

## Week 1: Project Setup & Design Specification

### Goals

- Finalize project scope and architecture.
- Prepare build environment and development workflow.
- Write pseudocode and class design.

### Milestones

#### 1. Literature & Design Research

- Review related work on distributed load balancing:
  - Chord / Consistent Hashing
  - Gossip protocols (e.g., SWIM, Serf)
  - Linux kernel scheduling as analogy
- Define how simulation models “load” (e.g., integer queue length).

#### 2. System Design Document

- Diagram the architecture:
  - PeerNode class (thread or process)
  - Message class for exchanging load and task info
  - NetworkSimulator or socket layer abstraction
- Define protocols (message types, update frequency).

#### 3. Environment Setup

- Initialize C++ project with CLion.
- Set up GitHub repo and build scripts.

- Verify multithreading, networking, and logging libraries (e.g., <thread>, <mutex>, <asio>).

#### 4. Deliverable:

- design\_spec.pdf including architecture diagram, pseudocode, and workflow.
- Confirm project compiles (“Hello World” test).

## Week 2: Core Implementation (Single-Node Simulation)

### Goals

- Implement node logic and basic local routing.
- Simulate load change and task queueing.

### Milestones

#### 1. Implement Node Core

- PeerNode manages:
  - A local load queue (tasks)
  - Thread-safe metrics for current load
  - Periodic load adjustment
- Add worker threads to “process” tasks (sleep-based simulation).

#### 2. Add Basic Task Routing

- Local node receives tasks.
- If load exceeds threshold, offload to random peer (placeholder routing).

#### 3. Logging and Metrics

- Log per-node load, task transfers, and timing.
- Implement local performance stats collection.

#### 4. Deliverable:

- Single-process simulation with 3–5 nodes (threads).
- Output showing dynamic task balancing locally.

## Week 3: Distributed Communication Layer

### Goals

- Implement inter-node communication.
- Replace placeholder routing with peer-to-peer message passing.

### Milestones

#### 1. Networking Layer

- Implement communication via:
  - **Option A:** TCP sockets (real networking)
  - **Option B:** Message queues or shared memory (simulated networking)
- Define message types: LOAD\_UPDATE, TASK\_REQUEST, TASK\_TRANSFER.

#### 2. Peer Discovery

- Implement static peer list or basic gossip broadcast.
- Peers periodically exchange load info.

#### 3. Decentralized Routing

- Each node uses peer load info to pick the least-loaded node.
- Route new or excess tasks dynamically.

#### 4. Deliverable:

- Multi-process distributed load balancer demo.
- Logs showing load distribution and routing stability.

## Week 4: Testing, Analysis, and Report Preparation

### Goals

- Evaluate performance, complete the short paper, and polish code.
- Connect findings to MoE research framework.

### Milestones

#### 1. Performance Evaluation

- Run tests varying number of peers and load intensity.
- Collect metrics: mean load variance, message rate, convergence time.

#### 2. Analysis

- Compare results to ideal load distribution.
- Identify bottlenecks or instability causes.

#### 3. Documentation

- Write **5 page report**:
  1. Introduction & Goal
  2. Related Work (Distributed scheduling, gossip routing)
  3. Interface (Sockets, threads, system calls used)
  4. Software Implementation (class structure, flow diagrams)
  5. Conclusion (difficulties, future work)
- Include architecture diagram and simulation graphs.

#### 4. Code Finalization

- Add README with compilation/run instructions.
- Comment all public classes and functions.

#### 5. Deliverables:

- Final C++ source code in GitHub repo.
- Completed 5 page project paper.
- Performance plots (optional).

### Stretch Goals

- Add a **Trust Score** to each node (randomly fluctuating or based on past reliability).
- Modify routing to avoid low-trust peers (basic trust-aware scheduling).
- Implement lightweight encryption or signature in messages (OpenSSL or libsodium).
- Visualize peer load over time using matplotlib or gnuplot.